

CSc 174

Database Management Systems

10. Relational Database Design Algorithms

Ying Jin

Computer Science Department

California state University, Sacramento



1. Properties of Relational Decompositions

Properties

- ◆ Attribute preservation
- ◆ Dependency preservation property
- ◆ Lossless join property

Universal Relation

◆ Universal Relation Schema:

a relation schema $R = \{A_1, A_2, \dots, A_n\}$ that includes **all** the attributes of the database.

◆ Universal relation assumption

every attribute name is unique

◆ Decomposition:

decompose **R** into **$D = \{R_1, R_2, \dots, R_m\}$** , by using the functional dependencies.

Attribute Preservation

◆ Attribute preservation

Each attribute in R will appear in **at least one** relation schema R_i in the decomposition, so that no attributes are “lost”.

Dependency Preservation Property

-Informally

◆ Informally:

Each function dependency $X \rightarrow Y$ specified in F , either

(1) **appeared directly** in one of the relation schemas R_i in the decomposition D , or

(2) could be **inferred** from the dependencies that appear in some R_i .

◆ What is F^+ (Closure of F)?

Project of F on R - Definition

◆ Definition:

Given a set of dependencies F on R , the **projection** of F on R_i , denoted by $\pi_{R_i}(F)$, where $R_i \subseteq R$, is the set of dependencies $X \rightarrow Y$ in F^+ such that the attributes in $X \cup Y$ are all contained in R_i .

all their left- and right-hand-side attributes are in R_i .

Dependency Preservation Property

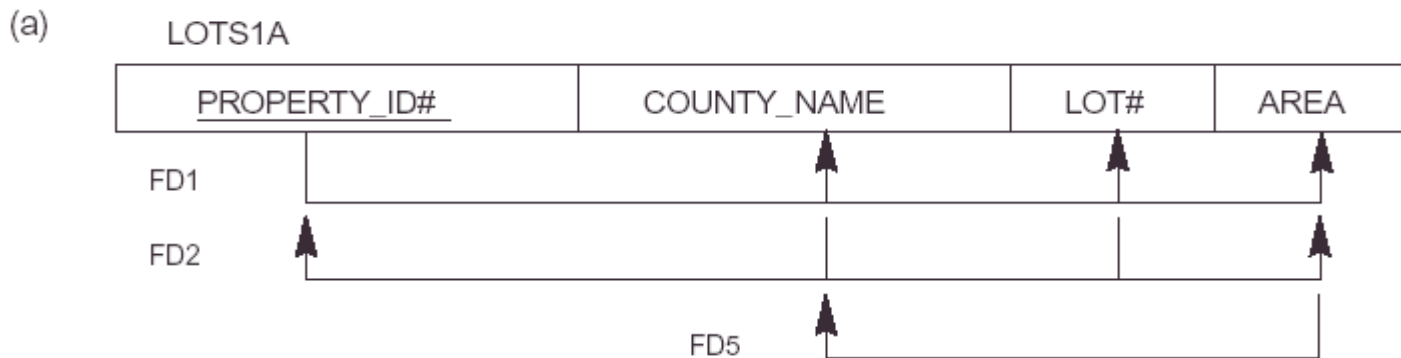
- ◆ Dependency Preservation Property definition:
a decomposition $D = \{R_1, R_2, \dots, R_m\}$ of R is **dependency-preserving** with respect to F if the union of the projections of F on each R_i in D is equivalent to F , that is, (

$$(\pi_{R_1}(F)) \cup \dots \cup (\pi_{R_m}(F))^+ = F^+$$

- ◆ **Claim 1:** It is always possible to find a **dependency-preserving** decomposition D with respect to F such that each relation R_i in D is in **3nf**.

Not Dependency-Preservation Example

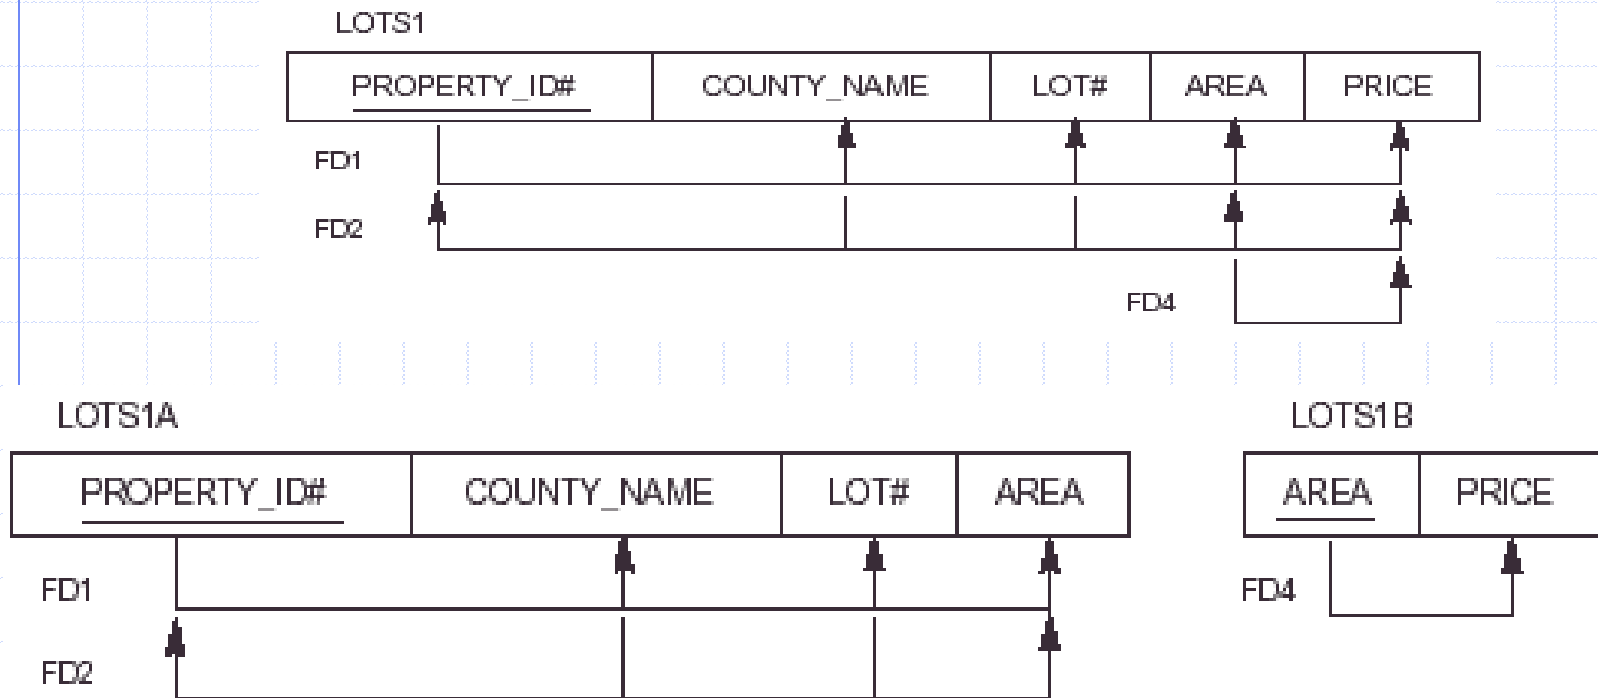
◆ Which fd is lost?



BCNF Normalization



Dependency-Preservation Example



Lossless (Non-additive) Join Property

- ◆ No **addition** of **spurious** tuples.

- ◆ Definition

Lossless join property: a decomposition $D = \{R_1, R_2, \dots, R_m\}$ of R has the **lossless (nonadditive) join property** with respect to the set of dependencies F on R if, for *every* relation state r of R that satisfies F , the following holds, where $*$ is the natural join of all the relations in D :

$$* (\pi_{R_1}(r), \dots, \pi_{R_m}(r)) = r$$

- ◆ **Note:** *lossless* refers to *loss of information*, not to loss of tuples.

Algorithm for Lossless (Non-additive) Join Property of a Decomposition

Testing for Lossless Join Property

Input: A universal relation R , a decomposition $D = \{R_1, R_2, \dots, R_m\}$ of R , and a set F of functional dependencies.

1. Create an initial matrix S with one row i for each relation R_i in D , and one column j for each attribute A_j in R .
2. Set $S(i,j) := b_{ij}$ for all matrix entries. (each b_{ij} is a distinct symbol associated with indices (i,j)).
3. For each row i representing relation schema R_i
 {for each column j representing attribute A_j
 {if (relation R_i includes attribute A_j) then set $S(i,j) := a_{ji}$ };};
(each a_j is a distinct symbol associated with index (j))

Algorithm for Lossless (Non-additive) Join Property of a Decomposition (Cont.)

4. Repeat the following loop until a *complete loop execution* results in no changes to S
 - {for each functional dependency $X \rightarrow Y$ in F
 - {for all rows in S which have the same symbols in the columns corresponding to attributes in X
 - {make the symbols in each column that correspond to an attribute in Y be the same in all these rows as follows: if any of the rows has an “ a ” symbol for the column, set the other rows to that *same* “ a ” symbol in the column; If no “ a ” symbol exists for the attribute in any of the row, choose one of the “ b ” symbols that appears in one of the rows for the attribute and set the other rows to that same “ b ” symbol in the column };
5. If a row is made up entirely of “ a ” symbols, then the decomposition has the lossless join property; otherwise it does not.

Lossless (nonadditive) join test for n -ary decompositions – Example

Decomposition of EMP_PROJ into EMP, PROJECT, and WORKS_ON satisfies test.

EMP

SSN	ENAME
-----	-------

PROJECT

PNUMBER	PNAME	PLOCATION
---------	-------	-----------

WORKS_ON

SSN	PNUMBER	HOURS
-----	---------	-------

(c)

$R = \{SSN, ENAME, PNUMBER, PNAME, PLOCATION, HOURS\}$

$D = \{R_1, R_2, R_3\}$

$R_1 = EMP = \{SSN, ENAME\}$

$R_2 = PROJ = \{PNUMBER, PNAME, PLOCATION\}$

$R_3 = WORKS_ON = \{SSN, PNUMBER, HOURS\}$

$F = \{SSN \rightarrow \{ENAME\}; PNUMBER \rightarrow \{PNAME, PLOCATION\}; \{SSN, PNUMBER\} \rightarrow HOURS\}$

	SSN	ENAME	PNUMBER	PNAME	PLOCATION	HOURS
R_1	a_1	a_2	b_{13}	b_{14}	b_{15}	b_{16}
R_2	b_{21}	b_{22}	a_3	a_4	a_5	b_{26}
R_3	a_1	b_{32}	a_3	b_{34}	b_{35}	a_6

(original matrix S at start of algorithm)

	SSN	ENAME	PNUMBER	PNAME	PLOCATION	HOURS
R_1	a_1	a_2	b_{13}	b_{14}	b_{15}	b_{16}
R_2	b_{21}	b_{22}	a_3	a_4	a_5	b_{26}
R_3	a_1	b_{32} a_2	a_3	b_{34} a_4	b_{35} a_5	a_6

(matrix S after applying the first two functional dependencies - last row is all "a" symbols, so we stop)

Binary Decomposition

- ◆ **Binary Decomposition:** decomposition of a relation R into two relations.
- ◆ **PROPERTY LJ1 (lossless join test for binary decompositions):** A decomposition $D = \{R_1, R_2\}$ of R has the lossless join property with respect to a set of functional dependencies F on R *if and only if* either
 - The f.d. $((R_1 \cap R_2) \rightarrow (R_1 - R_2))$ is in F^+ , or
 - The f.d. $((R_1 \cap R_2) \rightarrow (R_2 - R_1))$ is in F^+ .

Successive Lossless Join Decomposition

Claim 2 (Preservation of non-additivity in successive decompositions):

If a decomposition $D = \{R_1, R_2, \dots, R_m\}$ of R has the lossless (non-additive) join property with respect to a set of functional dependencies F on R , and if a decomposition $D_i = \{Q_1, Q_2, \dots, Q_k\}$ of R_i has the lossless (non-additive) join property with respect to the projection of F on R_i , then the decomposition $D_2 = \{R_1, R_2, \dots, R_{i-1}, Q_1, Q_2, \dots, Q_k, R_{i+1}, \dots, R_m\}$ of R has the non-additive join property with respect to F .



2. Algorithms for Relational Database Schema Design

Relational Synthesis into 3NF with Dependency Preservation (*Relational Synthesis Algorithm*)

◆ Does not guarantee the lossless join property

Input: A universal relation R and a set of functional dependencies F on the attributes of R .

1. Find a minimal cover G for F (use Algorithm 10.2);
2. For each left-hand-side X of a functional dependency that appears in G , create a relation schema in D with attributes $\{X \cup \{A_1\} \cup \{A_2\} \dots \cup \{A_k\}\}$, where $X \rightarrow A_1, X \rightarrow A_2, \dots, X \rightarrow A_k$ are the only dependencies in G with X as left-hand-side (X is the *key* of this relation) ;
3. Place any remaining attributes (that have not been placed in any relation) in a single relation schema to ensure the attribute preservation property.

Relational Decomposition into BCNF with Lossless (non-additive) join property

◆ No guarantee of dependency preservation

Input: A universal relation R and a set of functional dependencies F on the attributes of R .

1. Set $D := \{R\}$;
2. While (there is a relation schema Q in D that is not in BCNF) do {
 choose a relation schema Q in D that is not in BCNF;
 find a functional dependency $X \rightarrow Y$ in Q that violates BCNF;
 replace Q in D by two relation schemas $(Q - Y)$ and $(X \cup Y)$;
};

Relational Synthesis into 3NF with Dependency Preservation *and* Lossless (Non-Additive) Join Property

- ⊕ **Input:** A universal relation R and a set of functional dependencies F on the attributes of R .
1. Find a minimal cover G for F (*Use Algorithm 10.2*).
 2. For each left-hand-side X of a functional dependency that appears in G , create a relation schema in D with attributes $\{X \cup \{A_1\} \cup \{A_2\} \dots \cup \{A_k\}\}$, where $X \rightarrow A_1$, $X \rightarrow A_2$, ..., $X \rightarrow A_k$ are the only dependencies in G with X as left-hand-side (X is the *key* of this relation).
 3. If none of the relation schemas in D contains a key of R , then create one more relation schema in D that contains attributes that form a key of R . (*Use Algorithm 11.4a to find the key of R*).
 4. Eliminate redundant relations from the resulting set of relations in the relational database schema. A relation R is considered redundant if R is a project of another relation S in the schema.

Algorithm: Finding a Key K for R Given a set F of Functional Dependencies

Input: A universal relation R and a set of functional dependencies F on the attributes of R .

1. Set $K := R$.
2. For each attribute A in K {
 compute $(K - A)^+$ with respect to F ;
 If $(K - A)^+$ contains all the attributes in R ,
 then set $K := K - \{A\}$; }

Discussion of Normalization Algorithms

Problems:

- ◆ The database designer must first specify *all the relevant functional dependencies* among the database attributes.
- ◆ It is not always possible to find a decomposition into relation schemas that preserves dependencies and allows each relation schema in the decomposition to be in **BCNF** (instead of 3NF as in Algorithm 11.4).
- ◆ Many *different designs* arise corresponding to the same set of functional dependencies, depending on the *order* in which such dependencies are considered.
- ◆ Some design may be superior, or undesirable.

Algorithms for Relational Database Schema Design

	Input	Output	Properties/Purpose	Remarks
	A decomposition D of R and a set F of functional dependencies	Boolean result: yes or no for lossless join property	Testing for lossless join decomposition	
	Set of functional dependencies F	A set of relations in 3NF	Dependency preservation	No guarantee of satisfying lossless join property
	Set of functional dependencies F	A set of relations in BCNF	Lossless join decomposition	No guarantee of dependency preservation
	Set of functional dependencies F	A set of relations in 3NF	Lossless join and dependency preserving decomposition	May not achieve BCNF
	Relation schema R with a set of functional dependencies F	Key K of R	To find a key K (which is a subset of R)	The entire relation R is always a default superkey



These slides are based on the textbook:

R. Elmasri and S. Navathe, *Fundamentals of Database System*, 7th Edition, Addison-Wesley.