



CSc 174

Database Management Systems

13. Concurrency Control

Ying Jin

Computer Science Department

California state University, Sacramento

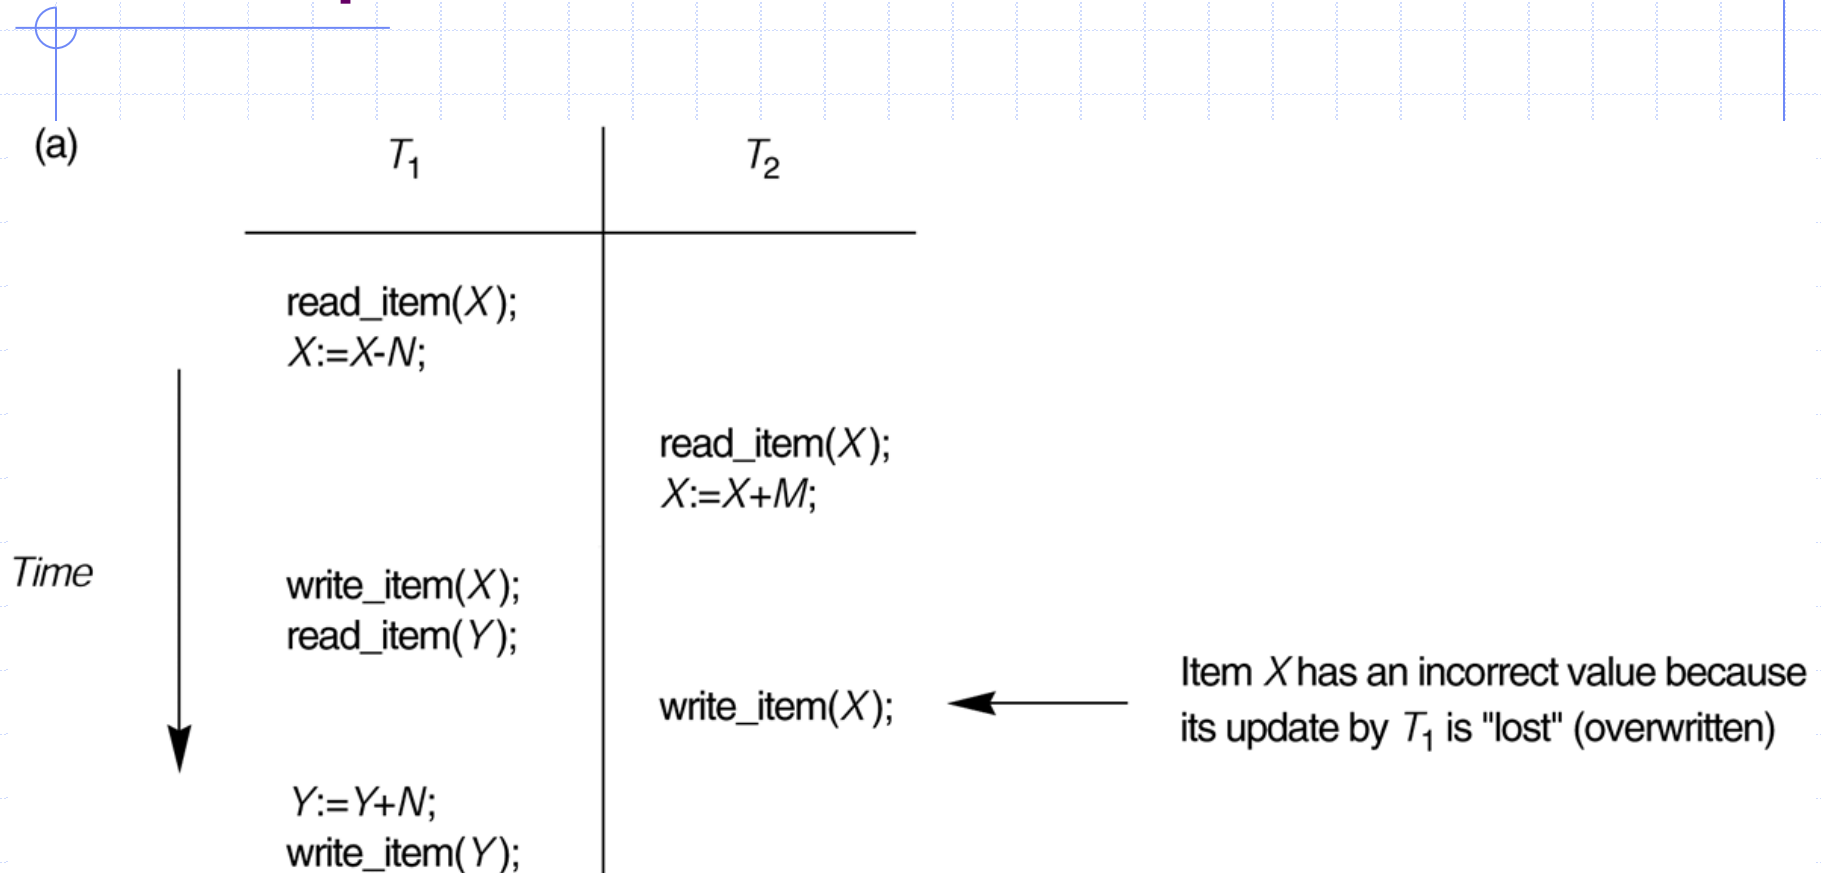
Concurrency Control

- ◆ Transactions can execute concurrently.
- ◆ Problems when allow concurrent execution without control
 - **The Lost Update Problem.**

This occurs when

- two transactions that **access the same** database items and
- have their operations **interleaved** in a way that makes the value of some database item incorrect.

The Lost Update Problem - Example



Dirty Read problem

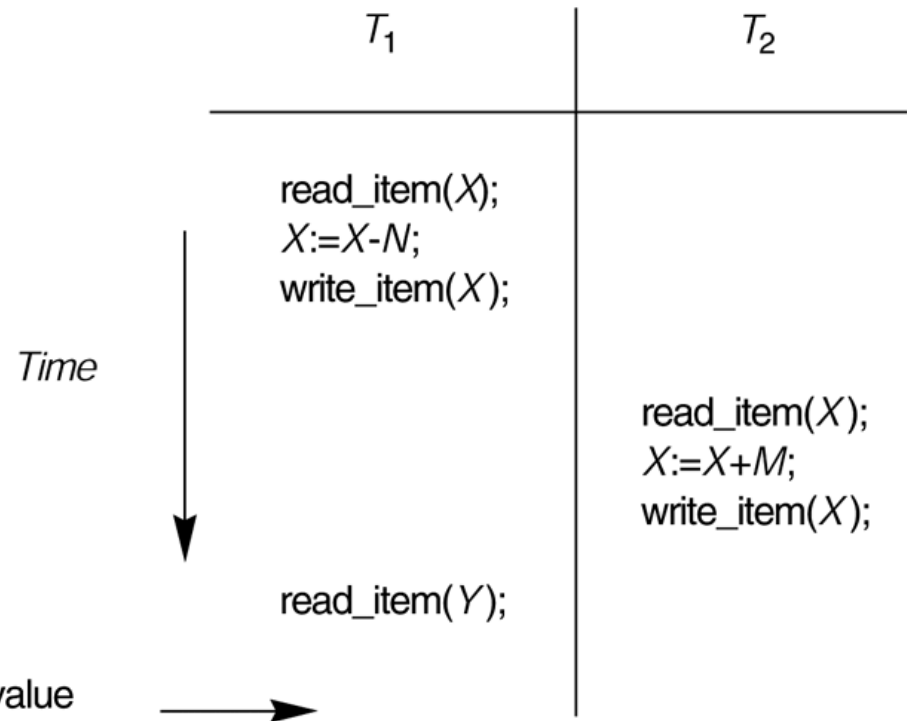
◆ The Temporary Update (or Dirty Read) Problem.

This occurs when

- one transaction updates a database item
- and then the **transaction fails** for some reason.
- The updated item is accessed by another transaction before it is changed back to its original value.

Dirty Read problem - Example

(b)



Transaction T_1 fails and must change the value of X back to its old value; meanwhile T_2 has read the "temporary" incorrect value of X .

Solve problems

- ◆ Problems occurs without control.
- ◆ Concurrency Control
 - Concepts: schedule, serializable
 - Concurrency control technique

Transaction Schedule

◆ A **schedule** (or **history**) S of n transactions T_1, T_2, \dots, T_n

It is an ordering of the operations of the transactions subject to the constraint that:

- for each transaction T_i that participates in S
- the operations of T_i in S must appear in the same order in which they occur in T_i .

◆ Note, however, that operations from other transactions T_j can be interleaved with the operations of T_i in S .

Conflict operations

◆ Two operations in a schedule **conflict** if they satisfy all three of the following conditions:

- (1) They belong to different transactions
- (2) They access the same item X
- (3) At least one of the operations is a write_item(X)

Equivalent schedule

- ◆ **Result equivalent:** Two schedules are called result equivalent if they produce the same final state of the database.
 - Two different schedules may accidentally produce the same final state.
- ◆ **Conflict equivalent:** Two schedules are said to be conflict equivalent if the order of any two conflicting operations is the same in both schedules.

Serial Schedule

- ◆ A schedule S is **serial** if,
 - for every transaction T participating in the schedule,
 - all the operations of T are executed consecutively in the schedule.

(Without interleaving - Only one transaction active at once)

- ◆ Otherwise, the schedule is called **nonserial schedule**.

Serializable Schedule

- ◆ **Conflict serializable:** A schedule S is said to be conflict serializable if it is conflict equivalent to some serial schedule S' .
- ◆ We call it “serializable” in the rest of this CSC 174 class.
- ◆ Being serializable is not the same as being serial

Property of Serializable schedule

- ◆ Why study Serializability?
- ◆ Being serializable implies that the schedule is a **correct** schedule.
 - It will leave the database in a consistent state.
 - The interleaving is appropriate and will result in a state **as if the transactions were serially executed**, yet will achieve efficiency due to concurrent execution.

Testing for conflict serializability

Algorithm 17.1:

1. Looks at only read_Item (X) and write_Item (X) operations
2. Constructs a precedence graph (serialization graph) - a graph with directed edges
3. An edge is created from T_i to T_j if one of the operations in T_i appears before a conflicting operation in T_j (RW,WR,WW)
4. The schedule is serializable if and only if the precedence graph has no cycles.

● Example 1

Example 2

(a)

transaction T_1

```
read_item (X);  
write_item (X);  
read_item (Y);  
write_item (Y);
```

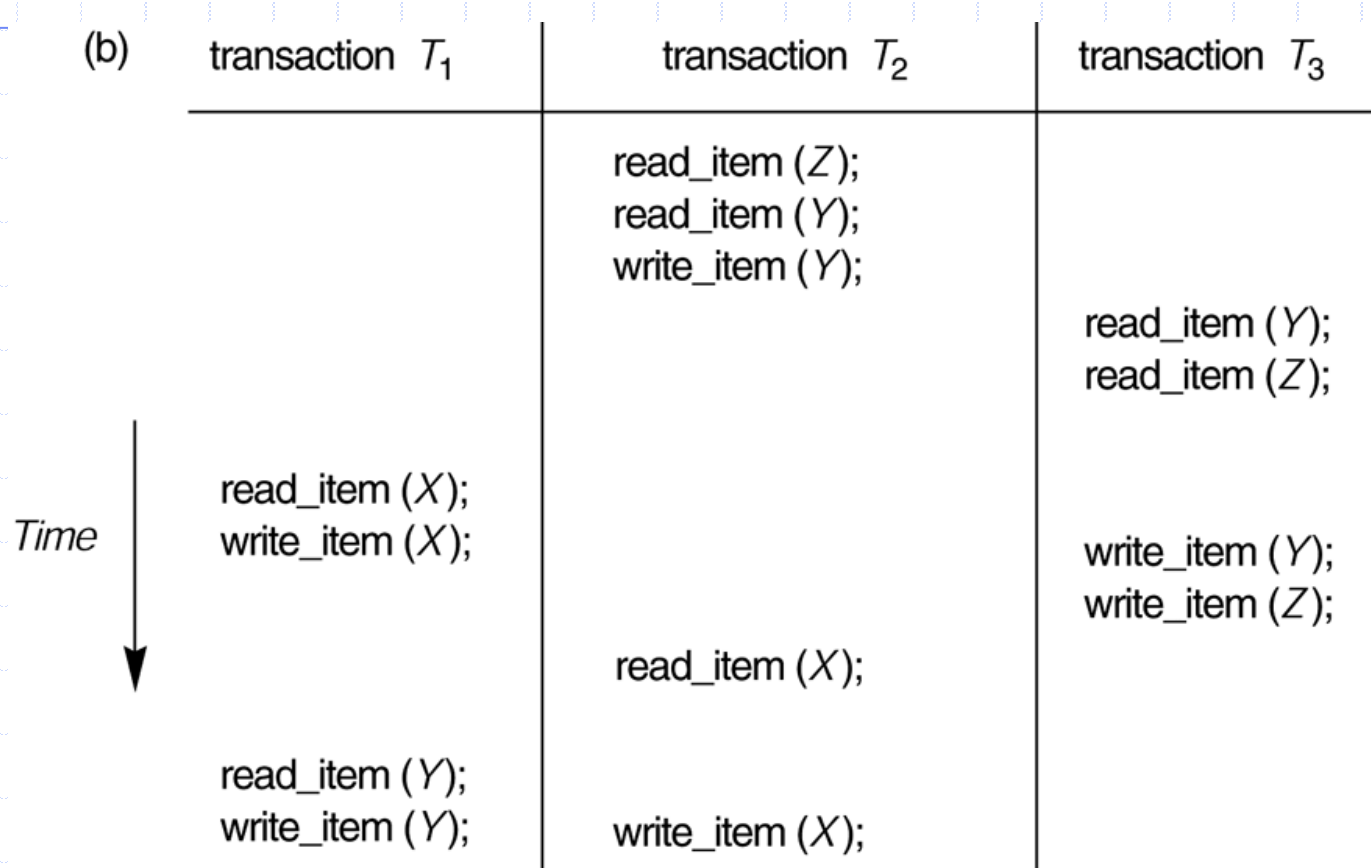
transaction T_2

```
read_item (Z);  
read_item (Y);  
write_item (Y);  
read_item (X);  
write_item (X);
```

transaction T_3

```
read_item (Y);  
read_item (Z);  
write_item (Y);  
write_item (Z);
```


Example 2 (cont.)



Schedule E

Example 2 (Cont.)

(d)



Equivalent serial schedules

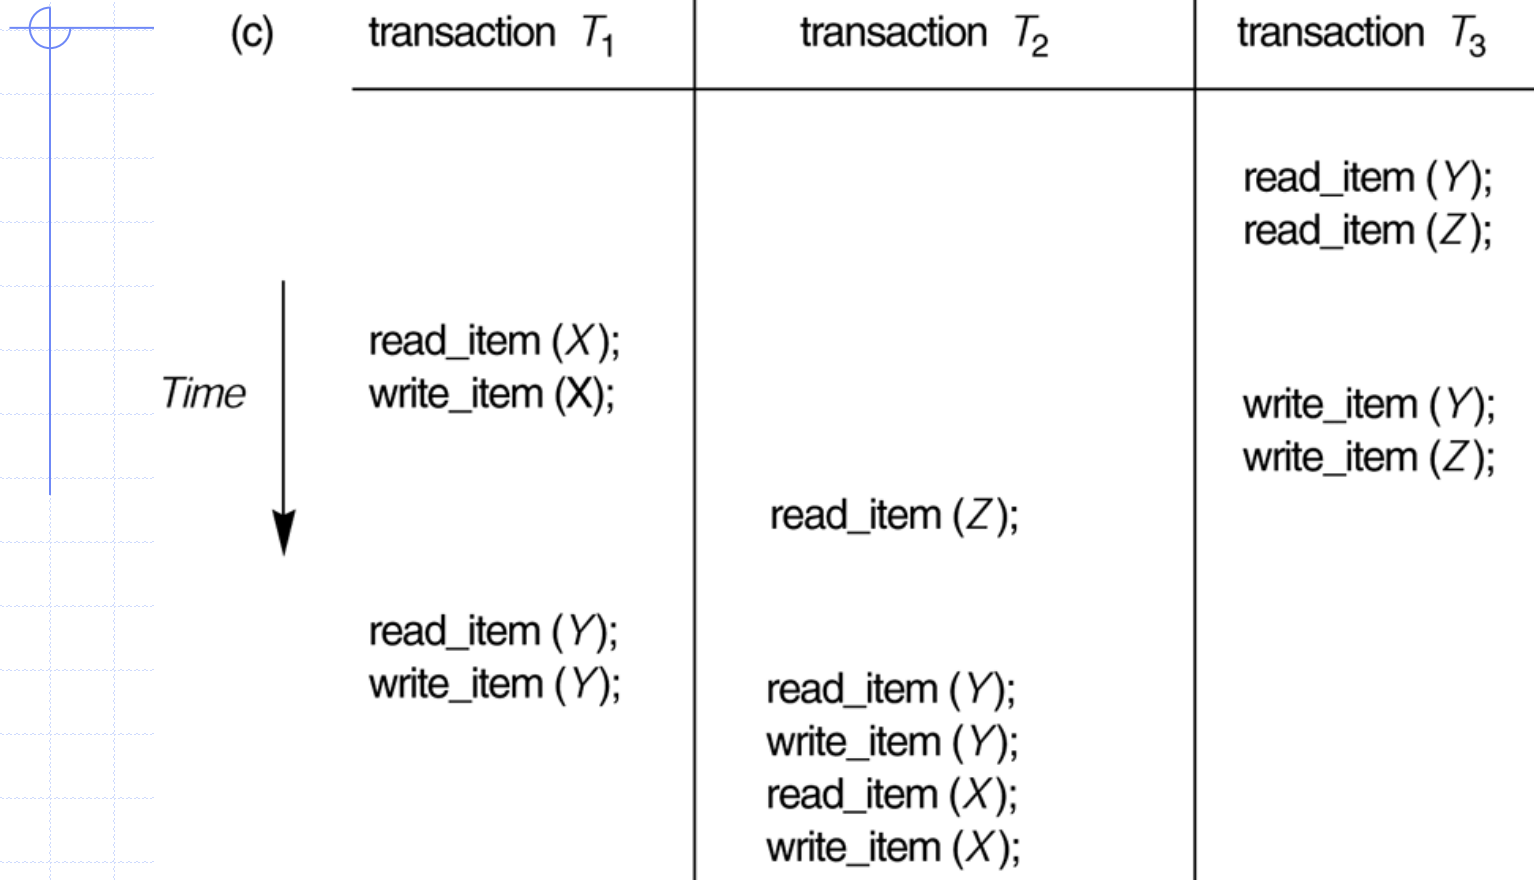
None

Reason

cycle $X(T_1 \rightarrow T_2), Y(T_2 \rightarrow T_1)$

cycle $X(T_1 \rightarrow T_2), YZ(T_2 \rightarrow T_3), Y(T_3 \rightarrow T_1)$

Exercise



Schedule F

Use of Serializability

◆ In practice

- Interleaving of operations from concurrent transactions is typically determined by the operating system scheduler
- It's difficult to determine beforehand how the operations in a schedule will be interleaved.

- ◆ If transactions are executed at will and then resulting schedule is tested for serializability,
 - cancel the effect of non-serializable schedule.
 - impractical

Practical Approach

- ◆ Come up with methods (protocols) to ensure serializability.
- ◆ Current approach used in most DBMSs:
Concurrency Control Techniques
 - Two-Phase locking
 - Timestamp Ordering

Lock

Locking is an operation which secures

- (a) permission to Read or
- (b) permission to Write a data item for a transaction.

Binary Locks:

- ◆ Lock (X). Data item X is locked in behalf of the requesting transaction.
- ◆ Unlock (X). Data item X is made available to all other transactions.

Rules for Binary locking

- ◆ Every transaction T must obey the following rules:
 - T must issue the operation lock(x) before any read(x) or write(x) operations are performed in T.
 - T must issue the operation unlock(x) after all read(x) and write(x) operations are completed in T.
 - T will not issue lock(x) operation if it already holds the lock on item x
 - T will not issue an unlock(x) operation unless it already holds the lock on item X.
- ◆ At most one transaction can hold the lock on a particular item.

Shared/Exclusive locks

- ◆ The above binary locking is too restrictive.
- ◆ We should allow several transactions to access the same item X if they all access x for **reading** purpose
- ◆ Three locking operations: `read_lock(x)`, `write_lock(x)`, and `unlock(x)`.

Shared lock: `read_lock (X)`. More than one transaction can apply share lock on X for reading its value.

Exclusive lock: `write_lock(x)`. Only a single transaction exclusively holds the lock on the item.

Lock Compatibility

- ◆ T is requesting a lock.
- ◆ Lock: The lock manager issue a lock to T.
- ◆ Wait: T has to wait until the lock is issued by the lock manager.

| | READ REQUEST | WRITE REQUEST |
|------------|--------------|---------------|
| READ LOCK | LOCK | WAIT |
| WRITE LOCK | WAIT | WAIT |

Well-formed transactions

- ◆ Database requires that all transactions should be well-formed.
- ◆ A transaction is well-formed if:
 - It must lock the data item in an **appropriate mode** before it reads or writes to it. (read: read_lock, write: write_lock)
 - If the data item is already locked by another transaction in an incompatible mode, then T **must wait** until all incompatible locks held by other transactions have been released.
 - It must **not** try to **unlock a free data item**.

Locking Example

T1

```
read_lock (Y);  
read_item (Y);  
unlock (Y);  
write_lock (X);  
read_item (X);  
X:=X+Y;  
write_item (X);  
unlock (X);
```

T2

```
read_lock (X);  
read_item (X);  
unlock (X);  
Write_lock (Y);  
read_item (Y);  
Y:=X+Y;  
write_item (Y);  
unlock (Y);
```

Result

Initial values: X=20; Y=30
Result of serial execution
T1 followed by T2
X=50, Y=80.
Result of serial execution
T2 followed by T1
X=70, Y=50

Locking Example (Cont)

T1

**read_lock (Y);
read_item (Y);**

unlock (Y);

**write_lock (X);
read_item (X);
X:=X+Y;
write_item (X);
unlock (X);**

T2

**read_lock (X);
read_item (X);
unlock (X);
write_lock (Y);
read_item (Y);
Y:=X+Y;
write_item (Y);
unlock (Y);**

Result

**X=50; Y=50
Well-formed.**

Serializable?

**Lock guarantee Serializable?
Need more rules/ protocols.**

Two-Phase Locking Protocol (2PL)

- ◆ A transaction is said to follow the **two-phase locking protocol** if all locking operations (read_lock, write_lock) precede the *first* unlock operation.
- ◆ Transaction can be divided into two phases
 - Growing (first) phase: new locks on items can be acquired but none can be released
 - Shrinking(second) phase: existing locks can be released but no new locks can be acquired.

2PL (Cont.)

- ◆ If every transaction in a schedule follows the two-phase locking protocol, the schedule is guaranteed to be serializable.

Example 1 of 2PL

T'1

```
read_lock (Y);  
read_item (Y);  
write_lock (X);  
unlock (Y);  
read_item (X);  
X:=X+Y;  
write_item (X);  
unlock (X);
```

Example 2 of 2PL

T'2

```
read_lock (X);  
read_item (X);  
Write_lock (Y);  
unlock (X);  
read_item (Y);  
Y:=X+Y;  
write_item (Y);  
unlock (Y);
```

◆ Example 3: Interleaving of T3 and T4.

Deadlock

- ◆ **Deadlock** occurs when each transaction T in a set of two or more transactions is waiting for some item that is locked by some other transaction T' in the set.

Deadlock Example

T'1

read_lock (Y);
read_item (Y);

write_lock (X);
(waits for X)

T'2

read_lock (X);
read_item (X);

write_lock (Y);
(waits for Y)

T1 and T2 did follow two-phase policy, but they are deadlock

Deadlock (T'1 and T'2)

Deadlock Prevention

- ◆ Use deadlock prevention protocol to prevent deadlock.
 - Conservative two-phase locking:
 - ◆ Requires a transaction to lock all the items it accesses before the transaction begins execution.
 - ◆ If any of the items cannot be obtained, none of the items are locked.
 - ◆ Then the transaction will try again at a later time.
 - Not practical


Deadlock Detection

- ◆ DBMS checks if a state of deadlock actually exists.
- ◆ More practical approach
- ◆ Wait-for graph
 - One node for each transaction that is currently executing
 - Whenever **T_i is waiting** to lock an item X that is currently locked by a transaction T_j, a directed edge (**T_i** -> T_j) is created in the wait-for graph

Wait-for graph example

T1' and T2' of Page 33

(b)



Deadlock Detection (Cont.)

- ◆ Deadlock if and only if the wait-for graph has cycle
- ◆ DBMS needs to select a time to check for a deadlock.

Deadlock Detection (Cont.)

- ◆ When a deadlock occurs, DBMS will choose some transaction causing the deadlock to abort.
- ◆ Victim Selection algorithm:
 - Avoid selecting transactions that have been running for a long time and that have performed many update.
 - Select transactions that have not made many changes

Starvation

- ◆ Starvation occurs when a transaction cannot proceed for an indefinite period of time while other transactions in the system continue normally.
- ◆ This may occur if the waiting/aborting scheme for locked items is unfair.
 - Fair waiting scheme: first-come-first-served
 - Allow priority: Priority increasing during waiting.

Timeout

- ◆ Simple approach to deal with deadlock
- ◆ If a transaction waits for a period longer than a system-defined timeout period, the system assumes that the transaction may be deadlocked and aborts it – **regardless of whether a deadlock actually exists or not.**
- ◆ Practical because of its low overhead and simplicity.



These slides are based on the textbook:

R. Elmasri and S. Navathe, *Fundamentals of Database System*, 7th Edition, Addison-Wesley.